

# Efficient vectors with static and dynamic storage in C++

Hubert Obrzut

Institute of Compute Science, University of Wrocław

Thursday 17<sup>th</sup> February, 2022

# Presentation plan

- 1 *vector* overview
- 2 Static storage variant
- 3 Dynamic storage variant
- 4 Overriding type triviality
- 5 Summary

# vector overview

## Characteristics

- simple purpose: storage for objects, managing their lifetime and random access
- the most commonly used data structure in C++
- several variants and implementations

# vector overview

## Characteristics

- simple purpose: storage for objects, managing their lifetime and random access
- the most commonly used data structure in C++
- several variants and implementations

**Goal:** efficient implementation of two variants: with static and dynamic storage.

# static\_vector data structure

```
1 template<class T, size_t C>  
2 class static_vector<T, C>;
```

Listing 1: static\_vector declaration

## Overview

- static storage
- dynamically-resizable

## Usage

Useful when:

- number of elements to store is known beforehand
- dynamic allocation is not acceptable, e.g. performance reasons
- objects cannot be default constructed (example below)

# Usage example

```
1 T array[N];  
2 for (int i = 0; i < N; ++i)  
3     array[i] = T(custom, arguments);  
4  
5 static_vector<T, N> v;  
6 for (int i = 0; i < N; ++i)  
7     v.emplace_back(custom, arguments);
```

In C-style array case, we unnecessary create default constructed objects.

# Existing implementations

## `static_vector` implementations

- Boost
- Folly
- EASTL
- `uwr::static_vector`

# Existing implementations

## `static_vector` implementations

- Boost
- Folly
- EASTL
- `uwr::static_vector`

## Note

No STL implementation.



# Alternative implementations

```
1 template<class T, size_t C>
2 class static_vector {
3     ...
4     size_type m_size;
5     ...
6 };
```

Listing 2: *size-based* variant.

```
1 template<class T, size_t C>
2 class static_vector {
3     ...
4     T* m_end;
5     ...
6 };
```

Listing 3: *pointer-based* variant.

## Alternative implementations cont.

### *size-based* variant advantages

- slightly faster in the average case
- more natural
- minimal size dispatch
- *trivially-relocatable*

```
1 // size-based variant
2 assert(sizeof(static_vector<char, 5>) == 6);
```

Listing 4: Occupies 6 bytes instead of 16.

# uwr::static\_vector advantages


## Improvements

- implemented as a separate type – optimized for static storage
- cache-friendly
- push\_back function optimization
- economical size

## Additional features

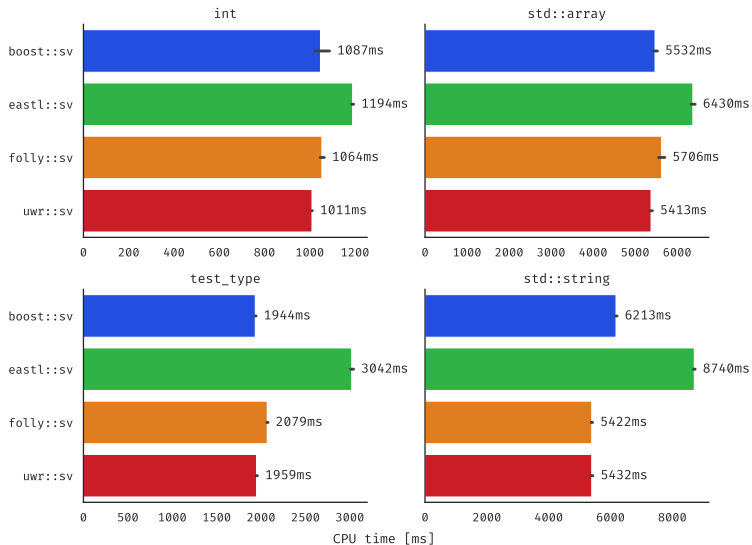
- C++20 compliant, C++17 compatible – operator<=, std::erase, std::erase\_if
- constexpr support
- fast\_push\_back, fast\_emplace\_back, safe\_pop\_back

# Benchmark



**Figure:** Benchmark example comparing `uwr::static_vector` with alternative implementations.

# Benchmark



## Real code example

### Chess benchmark

	Number of playouts	Number of states
<code>boost::static_vector</code>	625	195k
<code>uwr::static_vector</code>	765	240k

Aprox. 20% speedup.

# vector data structure

## Overview

- most commonly used variant of *vector*
- dynamically-resizable with dynamic storage
- keeps pointer to the allocated memory block
- continuous storage
- usually allocates more memory than needed
- uses growth policy for reallocation




Figure: *vector*'s growth example

# Existing implementations

## vector implementations

- STL
- Boost
- EASTL
- Folly
- rvector



# uwr::vector implementation

## Characteristics

- improves cache-friendliness
- uses memory mappings to allocate memory
- mmap, mremap, munmap

## mremap system call

- can expand mapping in-place
- very efficient for *trivial* types

```
1 return (T*)mremap(data,  
2                   old_capacity * sizeof(T),  
3                   new_capacity * sizeof(T),  
4                   MREMAP_MAYMOVE);
```

Listing 5: *trivial* type reallocation using mremap

## Reallocation strategy for *non-trivial* types

```
1 return (T*)mremap(data,  
2                   old_capacity * sizeof(T),  
3                   new_capacity * sizeof(T),  
4                   0);
```

Listing 6: *non-trivial* type reallocation using `mremap`

### Strategy

- `mremap` fails to expand in-place
- binary search the maximal successful growth over the number of pages

# Reallocation strategy for *non-trivial* types

```
1 return (T*)mremap(data,  
2                   old_capacity * sizeof(T),  
3                   new_capacity * sizeof(T),  
4                   0);
```

Listing 7: *non-trivial* type reallocation using `mremap`

## Strategy

- `mremap` fails to expand in-place
- binary search the maximal successful growth over the number of pages

## Results

Increased number of successful reallocations, e.g., from 24% to 45%.

# Benchmark

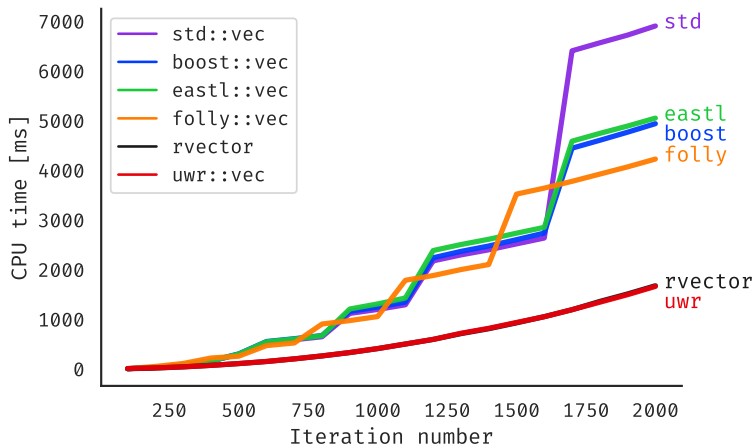


Figure: Example benchmark comparing `uwr::vector` with alternative implementations for *trivial* type.

# Benchmark

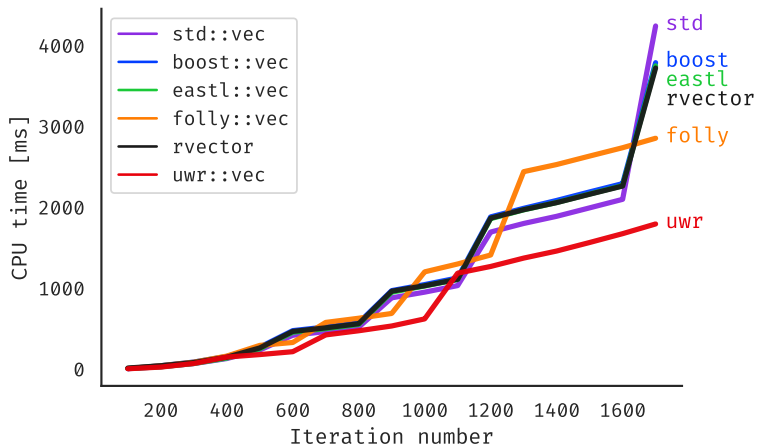


Figure: Example benchmark comparing `uwr::vector` with alternative implementations for *non-trivial* type.

# Benchmark




Figure: Example benchmark comparing `uwr::vector` with alternative implementations for *non-trivial* type with larger number of vectors in the environment.

# Custom growth rate

```
1 uwr::vector<T, std::ratio<7, 5>> v;
```

[Listing 8](#): Using `uwr::vector` with a custom growth rate.

# Type triviality

## *trivially-relocatable* type

Type that holds all of its invariants when moved to another memory location with raw copy.

## Note

Still can have defined move constructor/operator.



# Examples

## *trivial* types

- `uwr::vector`, `std::vector`, ...

# Examples

## *trivial* types

- `uwr::vector`, `std::vector`, ...
- `std::unordered_map`
- `std::unordered_set`

# Examples

## *trivial* types

- `uwr::vector`, `std::vector`, ...
- `std::unordered_map`
- `std::unordered_set`
- `std::unique_ptr`, `std::shared_ptr`

# Examples

## *trivial* types

- `uwr::vector`, `std::vector`, ...
- `std::unordered_map`
- `std::unordered_set`
- `std::unique_ptr`, `std::shared_ptr`
- big `std::string` (without SSO)

# Examples

## *trivial* types

- `uwr::vector`, `std::vector`, ...
- `std::unordered_map`
- `std::unordered_set`
- `std::unique_ptr`, `std::shared_ptr`
- big `std::string` (without SSO)

## Note

`std::map` and `std::set` are not.

# Overriding type triviality

```
1 namespace uwr::mem {  
2  
3 template<>  
4 inline constexpr bool is_trivially_relocatable_v<  
    MySpecialType> = true;  
5  
6 }
```

Listing 9: Overriding type triviality for `uwr::vector`

# Overriding type triviality

```
1 namespace uwr::mem {  
2  
3 template<>  
4 inline constexpr bool is_trivially_relocatable_v<  
    MySpecialType> = true;  
5  
6 }
```

Listing 10: Overriding type triviality for `uwr::vector`

## Note

`folly::vector` supports this feature as well.

# Benchmark




Figure: Example benchmark comparing `uwr::vector` with alternative implementations for type with overridden triviality.



# Summary

## Both vectors

- C++20 compliant, C++17 compatible
- cache-friendly
- no external dependencies

## `uwr::static_vector`

- economical size
- has `push_back` optimization

## `uwr::vector`

- allows overriding type triviality
- has custom *growth factor* support
- efficient for *trivial* and *non-trivial* types

# Thank you for your attention

## Both vectors

- C++20 compliant, C++17 compatible
- cache-friendly
- no external dependencies

## `uwr::static_vector`

- economical size
- has `push_back` optimization

## `uwr::vector`

- allows overriding type triviality
- has custom *growth factor* support
- efficient for *trivial* and *non-trivial* types